Attention-Based Deep Multi-Instance Learning for Weakly-labeled Breast Ultrasound Image Classification

Zena Fantaye

MSDS, University of Wisconsin – La Crosse

DS785: Capstone Project

Dr. Tracy Bibelnieks

Date: December 10, 2023

Abstract

This paper proposes using Attention-Based Multiple Instance Learning (ABMIL) on ultrasound images for breast cancer diagnosis. Rooted in an extensive literature review and building upon Shen et al.'s (2021) work, the study addresses challenges in traditional screening methods, particularly for cases with dense breast tissue. The paper provides a comprehensive overview of the research goals, literature review, methodology, and empirical findings. It thoroughly explores each aspect, giving insights into the details of the study.

The study compiles and processes 40,000 ultrasound images from the Mayo Clinic using a selective methodology to enhance diagnostic accuracy. The ABMIL model produced promising results, achieving an AUC-ROC of 0.8021, sensitivity of 0.73, and specificity of 0.87. The model's effectiveness with weakly labeled datasets is underscored by training and validation accuracy rates of 0.80 and 0.82, respectively. The model's interpretability is prioritized through attention mechanisms, corroborated by successful model application to an external dataset (Imagenette).

The research contributes to the evolving medical imaging landscape, exploring ABMIL's potential to advance breast cancer diagnosis and improve patient care.

Acknowledgments

I would like to express my sincere gratitude to Dr. Jeff Baggett from the University of Wisconsin-La Crosse for allowing me to contribute to this challenging yet rewarding project. His guidance, expertise, and framework have made this undertaking possible.

I want to extend special thanks to Dr. Rich Ellis and the Mayo Clinic for their generosity in providing the crucial dataset that forms the foundation of this research. I also want to thank Tristan Hansen for his significant contributions, particularly in data preparation and script development, which were crucial elements in the success of this study.

The collaborative efforts of these individuals and institutions were instrumental in bringing this research to fruition.

Table of Contents

Abstract....ii

Acknowledgments....iii

List of Tables....vi

List of Figures....vii

Chapter 1: Introduction....1

Background....1

The overarching project has three key objectives....2

Background of the Problem....2

Statement of the Problem....3

Purpose of the Study....4

Research Questions....4

Significance of the Study....5

Definition of Terms....5

Convolutional Neural network....5

Multiple instance learning (MIL)....5

Attention-based MIL Pooling (AbMILP)....6

Self-Attention Attention-based MIL Pooling....5

Assumption and Limitations....6

Conclusion....6

Chapter 2: LiteratureReview....8

Introduction....8

AI Role in Breast Ultrasound Imaging....8

Deep Multi-Instance Learning....9

Attention and Self-Attention Mechanism....10

Conclusion....12

Chapter 3: Methodology....14
Introduction....14
Data Collection....14
Data Processing....15
Model Architecture....16
Saliency Maps....17
Convolutional Layer....18
Sigmoid Activation....18
Flattening and Top-k Operation....18
Instance-Level Prediction....18
Attention Scores....18
Max-Pooling Operations....19
Linear Transformations....19
Attention Score....19
Aggregation and Bag-Level Prediction....19
Training and Evaluation....20
Summary....21
Chapter 4: Results....23
Introduction....23
Broad Results Overview....23
Effectiveness of the Attention-Based Multiple Instance Learning (ABMIL) Model....24
Interpretability of the ABMIL Model....28
Summary....31
Chapter 5: Discusion....33
Introduction....33
Summary of Findings....33

Conclusions (organized by project objectives) …….....34

Breast Cancer Diagnosis Accuracy Enhancement.....34

Interpretability Enhancement.....34

Generalization and Robust Pattern Recognition.....35

Discussion of Implications for Business.....35

Limitations of Results and Suggestions for Future Research.....36

Computational Challenges.....36

Correlation Assumption in the Model.....36

Challenges with Pretrained Encoder.....37

Dataset Specificity.....37

Conclusion.....38

References.....40

Appendix.....45

List of Tables

Table 1 Diagnostic Performance Metrics.....20

Table 2 Imaginate Dataset - Metrics.....29

List of Figures

Figure 1 Deep neural network architecture.....20

Figure 2 Training and Validation Losses Over 35 Epochs.....29

Figure 3 Validation Confusion Matrix.....20

Figure 4 Saliency Map – Highlighting Fish Among Images.....29

Figure 5 Example Saliency Maps from Research Conducted by Shen et al. (2021) .....,.29.

Chapter 1: Introduction

Breast cancer is the most prevalent form of cancer in adults, with over 2.3 million cases being diagnosed annually worldwide, according to the World Health Organization. In 95% of countries, it is the leading or second-leading cause of female cancer-related deaths (WHO, 2023). Detecting breast cancer early, before it metastasizes, is critical as it enables more effective treatments and improves survival rates (Shen et al., 2021). Although mammography is widely used for early breast cancer detection, it has some limitations, particularly for women with dense breast tissue. The sensitivity of mammography drops from 85% to 48-64% in such cases (Shen et al., 2021).

Due to mammography's limitations, ultrasound (US) plays a vital role as a supplementary modality in breast cancer screening and diagnosis. It is also one of the most preferred modalities of breast cancer detection due to its relatively low cost and ease of access.

Breast ultrasound has numerous benefits, but interpreting the results can be difficult as it heavily relies on the radiologist's judgment and experience. A radiologist evaluates lesions' size, shape, and orientation when examining the images to determine whether the findings are benign or malignant or need a suspicion-based biopsy (Shen et al., 2021). However, there is considerable variability in the recommendations made by radiologists, which can lead to more false-positive results and criticism of Breast US (Shen et al., 2021).

In the initial months of 2021, a collaborative initiative between the Mayo Clinic Health System and the University of Wisconsin-La Crosse was launched, focusing on developing computer-aided diagnosis (CAD) software. This software utilizes recent advancements in deep learning technology to enhance accuracy in breast cancer diagnosis. The project engages a diverse team comprising radiologists, mathematicians, data scientists, and graduate students, collectively

working towards creating software that rivals human experts in the precise interpretation and clinical assessment of breast lesions identified via ultrasound.

The overarching project has three key objectives:

Objective 1: Develop and train diverse neural network deep learning/machine models designed to classify breast ultrasound lesions.

Objective 2: Construct a secondary algorithm emulating systems employed by experts to ascertain a BI-RADS classification of a lesion. This system relies on key lesion characteristics, including shape, orientation, margins, internal and external features, elasticity, and vascularity – factors typically determined by radiologists.

Objective 3: Synthesize the two lesion assessment systems to yield an enhanced estimate of the BI-RADS Assessment Category.

This specific project contributes directly to Objective 1 of the larger initiative. Its primary goal is to improve the accuracy of breast cancer diagnosis in ultrasound images by developing and implementing a deep multi-instance learning (DML) model. The emphasis lies in addressing challenges presented by weakly labeled medical images, a domain where conventional machine learning approaches, reliant on fully annotated datasets, encounter limitations. The proposed DML model incorporates attention-based techniques to augment interpretability, ensuring transparency in decision-making processes for healthcare professionals. Through strategic utilization of a substantial dataset from the Mayo Clinic, the project aspires to make a meaningful contribution to the field by establishing a resilient and efficient system for breast cancer diagnosis.

Background of the Problem

Recent advances in AI/machine learning research, particularly convolutional neural networks (CNN) and related techniques, have boosted performance in image classification and

object detection tasks to be comparable to and even better than humans. Meanwhile, improvements in computing have significantly reduced the resources and time needed to process large amounts of data. These advances enable promising new approaches to building CAD systems.

Despite these advantages, interpreting breast ultrasound is challenging as it relies on a radiologist's judgment and experience. First, obtaining the large amount of clinical data required to train the neural network is difficult due to ownership and privacy issues. Second, the algorithm's performance largely depends on the data quality, where human expertise is needed to label a large set of images with high accuracy. Meanwhile, the structure of the neural network and the algorithm hyperparameters vary considerably for different organs, imaging modalities, and data sets, so there is no one-size-fits-all deep learning algorithm for CAD.

Statement of the Problem

From the beginning of the project, undergraduate and graduate students have been working with experts to build and train various deep learning/machine models to classify breast ultrasound lesions. Although significant progress has been made in building different algorithms, the proposed diagnostic approaches for breast cancer in ultrasound images face challenges related to the size and labeling of datasets, making it difficult to develop robust machine-learning models. Traditional algorithms that rely on well-labeled datasets struggle to adapt to the complexities of weakly labeled medical images. Additionally, many deep learning models lack interpretability, which poses challenges for healthcare professionals in understanding and trusting the decisions made by these systems. Therefore, this study aims to address these issues by proposing a DML model that incorporates attention-based mechanisms to improve the system's accuracy and interpretability.

Purpose of the Study

The main project began over two years ago to create software that can interpret and clinically assess breast lesions detected through ultrasound with the same level of accuracy as human experts. The current project aims to contribute to this main project by developing a more accurate and interpretable diagnostic DML model, overcoming the limitations of traditional machine-learning approaches. By leveraging attention-based techniques, the research seeks to optimize weakly labeled datasets, enhance interpretability for healthcare professionals, and adhere to ethical standards in handling patient data.

Research Questions

The proposed Attention-Based Multiple Instance Learning (ABMIL) model was studied to answer three key questions. Firstly, the research aimed to determine how effective the ABMIL model is in improving the accuracy of breast cancer diagnosis in ultrasound images. The primary research question focused on assessing the model's ability to differentiate between benign and malignant cases, emphasizing reducing false positives and unnecessary biopsies.

The second research question focused on the interpretability of the Deep Multi-Instance Learning (DML) model. The study examined how incorporating attention mechanisms and saliency maps enriched the model's interpretive capabilities. This question aimed to understand how these attention-based techniques provided insights into decision-making, offering a transparent view of the model's classifications.

Lastly, the research evaluated the performance of the ABMIL model in generalization and robust pattern recognition. The third question assessed how well the model adapted to previously unseen data, examining patterns in training and validation losses over epochs. The emphasis was

on uncovering the model's learning trajectory and identifying opportunities for optimization to enhance its overall performance in real-world scenarios.

Significance of the Study

The study that was conducted has significant implications for the diagnosis of breast cancer. Developing a precise and interpretable Deep Multiple Instance Learning (DML) model can revolutionize diagnostic methods by reducing the time and costs associated with the laborious task of labeling extensive datasets. The need for an accurate and automated system is particularly pronounced because ultrasound represents the clinical imaging modality most dependent on user interpretation and exhibits the widest range in radiologist lesion interpretation and positive biopsy rates. Internal studies at Mayo Clinic between January 1, 2019, and November 1, 2020, revealed a considerable fluctuation, ranging from 31% to 51%, in the percentage of positive breast ultrasound biopsies, indicative of non-uniform patient management. Deploying an automated system, especially for indeterminate lesions that evade easy classification as benign or malignant, can enhance patient care and concurrently alleviate medical costs.

Definition of Terms

***Convolutional Neural network*:** A neural network learns feature engineering using filter optimization.

Multiple instance learning (MIL): It is a weakly supervised learning algorithm where training data is arranged in sets called bags. Each bag contains a set of instances or images. Labels are provided for the entire sets or bags rather than for each image or instance inside the bag ("Improvising Weakly Supervised Object Detection (WSOD) Using Deep Learning Technique," 2020).

Attention-based MIL Pooling (AbMILP): It is a trainable operator that combines information from numerous instances within a bag. It is a two-layered neural network that leverages attention weights to identify essential instances (Rymarczyk et al., n.d.).

***Self-Attention Attention-based MIL Pooling*:** It is capable of simultaneously capturing global dependencies among instances in the bag and aggregating them into a fixed-sized vector needed for the subsequent layers of the network. This fixed-sized vector can be applied in regression, binary, and multi-class classification problems (Rymarczyk et al., n.d.).

Assumptions and Limitations

Given the extensive scope of the overarching project, distinct responsibilities have been assigned to different research team members. The focus of this study was confined to Objective 1, specifically the development of a deep-learning neural network model for classifying breast ultrasound lesions. It is assumed that the quality and accuracy of the Mayo Clinic dataset mirror those of more expansive breast cancer ultrasound image datasets. However, it is acknowledged that the study's reliance on the Mayo Clinic dataset may limit its generalizability to other datasets, contingent on the availability and quality of comparable data.

Conclusion

In summary, the collaborative effort initiated in early 2021 by the Mayo Clinic Health System and the University of Wisconsin-La Crosse represents a critical response to the persistent challenges in breast cancer diagnosis. As the most prevalent cancer globally, early detection is paramount, and the limitations of current methods, especially in ultrasound interpretation, necessitate innovative solutions. The project's focus on developing advanced computer-aided diagnosis (CAD) software, propelled by deep learning technologies, holds promise for overcoming the limitations of current diagnostic tools. By aligning with Objective 1 of the overarching project,

this research aims to contribute a deep multi-instance learning (DML) model, addressing issues associated with weakly labeled medical images and emphasizing interpretability. The anticipated outcomes carry the potential to reshape diagnostic practices, emphasizing efficiency, accuracy, and ethical considerations, thereby advancing the landscape of breast cancer diagnosis.

Chapter 2: Literature Review

The field of artificial intelligence (AI) has experienced incredible progress with increasing influence throughout the healthcare industry. Advanced AI analytics, such as deep learning, are playing a significant role in imagining analysis and classification in the medical field. Mayo Clinic, a nonprofit academic medical center, and the University of Wisconsin-La Crosse have been working together to develop computer-aided diagnosis (CAD) software that uses recent advances in deep learning technology to achieve better accuracy in breast ultrasound image classification. Currently, they are developing a deep multi-instance Learning model that can handle large, weakly labeled medical images and model dependencies between them while also being interpretable. To gain a better understanding of advancements in deep learning for image classification, specifically deep multi-instance Learning with attention and self-attention techniques, previous research was reviewed. Using Google Scholar, academic journal articles and research papers on breast ultrasound image classification with multiple instance learning methods were analyzed. This allowed for a comprehensive understanding of the current state of research in these domains and the identification of best practices.

AI Role in Breast Ultrasound Imaging

Breast cancer screening through mammography and ultrasound has limitations in detecting breast cancer, especially in women with dense breast tissue (Shen et al., 2021). Shen et al. (2021) conducted a study on the impact of artificial intelligence (AI) in reducing false-positive results in breast ultrasound exams. According to their research, mammography's sensitivity decreases from 85% to 48-64% in women with dense breast tissue (Shen et al., 2021). Moreover, ultrasound screenings produce more false-positive results, leading to 5-15% of patients being called back for further testing and 4-8% undergoing biopsy (Shen et al., 2021).

Such false-positive recalls and workups cost approximately $4 billion annually in the USA (Mayo et al., 2019). These false positives can cause unnecessary patient anxiety and may result in an unnecessary biopsy. This can lead to significant financial costs. Therefore, there is an urgent need for an AI system to improve the accuracy and efficiency of breast ultrasound analysis.

In their research, Shen et al. (2021) conducted a study where they introduced a deep learning model (DLM) capable of identifying breast cancers in ultrasound images with an accuracy level equivalent to that of a radiologist. Interestingly, the performance of a radiologist further improved when the deep learning model was used in combination with their expertise, demonstrating the effectiveness of a hybrid system. According to their research, with the help of the DLM, radiologists can decrease false positives by 37.3%, reduce biopsy requests by 27.8%, and maintain sensitivity (Shen et al., 2021). This is an encouraging start for many medical institutes, such as Mayo Clinic, who want to cut unnecessary medical procedures and medical costs. According to an internal study conducted at the Mayo Clinic, the rate of positive biopsy results following a breast ultrasound varied between 31% and 51% from 2019 to 2020. This indicates that many patients are subjected to unnecessary medical procedures and receive inconsistent standards of care, ultimately resulting in higher medical expenses. It also further strengthens the need for an AI system to improve breast ultrasound analysis's accuracy and efficiency.

Deep Multi-Instance Learning

Various Deep-learning algorithms are available for the analysis and classification of breast ultrasound. Among them, deep multi-Instance Learning has demonstrated encouraging outcomes in the medical field and has proven particularly effective in the context of medical

imaging. It is a weakly supervised learning algorithm where training data is arranged in sets called bags. Each bag contains a set of instances or images. Labels are provided for the entire sets or bags rather than for each image or instance inside the bag *("Improvising Weakly Supervised Object Detection (WSOD) Using Deep Learning Technique," 2020)*. Multiple Instance Learning (MIL) has two main approaches: Instance-based and Embedded/bag-based. The Instance-based approach involves learning an instance classifier and using its outputs to classify bags (Wang et al., 2018). However, this approach has some limitations as the instances inside a bag are not labeled, which leads to a lot of noise. Some researchers prefer the Embedded-based (bag-level) approach to overcome this issue. This approach involves embedding a bag into a vector representation, which is then used to apply classifiers for bag classification (Wang et al., 2018).

Wang et al. (2018b) propose a neural network framework for Multiple Instance Learning (MIL) in their research titled "Revisiting Multiple Instance Neural Networks." Their approach involves an embedded-based method that outperforms instance-based networks. Unlike instance-based networks, which use fully connected and activation layers to process each instance in a bag and obtain instance probabilities to determine bag probability, the proposed multi-instance neural network focuses on bag representation. This method classifies bags directly without calculating instance probability. During network training, the instances are initially labeled as latent variables. MIL pooling is then used to obtain a bag label (Wang et al., 2018b). However, the downside of this method is that it lacks interpretability. It does not identify the contribution of each instance to the bag label, making it challenging to identify the image with malignant lesions.

Attention and Self-Attention Mechanisms

Researchers have proposed incorporating an attention-based MIL pooling method to address the interpretability issues associated with the embedding-based approach. The main idea behind attention-based MIL pooling is to provide a more interpretable model for medical imaging tasks by better understanding the contribution of each instance to the bag label, unlike the MIL pooling method in the embedded approach. In the paper "Attention-based Deep Multiple Instance Learning," Ilse (2018) proposes a neural network-based permutation-invariant aggregation operator that corresponds to the attention mechanism. Their proposed attention-based method uses a weighted average of instances where a neural network determines weights. This method provides insight into the contribution of each instance to the bag label, with high attention weights assigned to instances that are likely to have a higher impact on the bag label (Ilse, 2018). However, their proposed approach is limited in assuming that all instances inside a bag are independent; it does not account for the dependency between instances inside a bag.

Similarly, Shen et al. (2021) propose a deep-learning neural network (DLM) that uses attention-based MIL pooling. The network first extracts features for each instance or image from a bag using the image-level information extractor block of the Convolutional Neural Network (CNN). The image-level feature maps then go through a process of one layer of convolution and sigmoid to produce saliency maps. These saliency maps highlight potentially benign and malignant lesions, making the model interpretable. The network also calculates attention scores that indicate the importance of the instances inside a bag for diagnosing benign and malignant lesions. Lastly, an information aggregator combines classification signals from all images to provide a breast-level prediction (Shen et al., 2021). The study shows that combining the radiologist's knowledge with the proposed deep learning model results in better performance and a more compelling hybrid system that reduces the number of false-positive findings in

interpreting breast ultrasound exams (Shen et al., 2021). However, the proposed model has limitations; although it aggregates information from the instances, it does not account for the dependency between instances inside a bag, similar to the previously mentioned attention-based neural network proposed by Ilse (2018).

Rymarczyk et al. (2021) recommend a new mechanism that combines self-attention with attention-based MIL Pooling to capture dependencies among instances in a bag. This method efficiently captures global dependencies among instances and generates a fixed-size vector for regression, binary, and multi-class classification tasks. The resulting vector is then passed to subsequent network layers, improving performance (Rymarczyk et al., 2021). The proposed method, Self-Attention into Attention-Based MIL Pooling (SAAbMILP), consists of four steps. Firstly, the images of the bag are sent through a Convolutional Neural Network (CNN) to obtain their features. The self-attention module then utilizes these features, which leverage dot product or other kernels to consider dependencies between the instances. The resulting feature vectors, which have integrated dependencies, are then passed through the AbMILP module to generate one fixed-sized vector per bag. This vector can be further fed into the successive Fully Connected (FC) network layers for classification (Rymarczyk et al., 2021).

Compared to the previous attention based DLM, this method includes a self-attention layer in the aggregator part of the model to capture the correlation between instances in a bag. However, unlike the attention based DLM, this method skips the process of passing through one layer of convolution and sigmoid to produce saliency maps. As a result, the method becomes less interpretable.

Conclusion

In the medical field, the use of AI systems is becoming increasingly necessary to improve the accuracy and efficiency of breast ultrasound analysis. When combined with a radiologist's expertise, these systems can result in a highly effective hybrid model. AI systems can significantly improve the analysis and classification of medical images. However, developing an Attention-Based deep multi-instance learning model that can handle large, weakly labeled medical images is crucial. Such a system should also be easy to interpret. While other researchers have developed similar models, the models developed and trained in the collaboration project between the Mayo Clinic and the University of Wisconsin since 2021 have been traditional machine learning models that use labeled datasets. Developing an interpretable model that can handle the vast, unlabeled, high-quality data available at the Mayo Clinic is essential.

Chapter 3: Methodology

The research methodology aligned to enhance breast cancer diagnosis accuracy in ultrasound images. The proposed approach involved creating a deep multi-instance learning model using attention-based techniques to deal with the challenges of medical image analysis. One of the main challenges is that many medical images were weakly labeled, and as mentioned previously, annotating these images is often both expensive and time-consuming. Rather than relying on labeled datasets like most traditional machine learning algorithms do, the aim was to develop a system that could effectively utilize the large number of unlabeled or weakly labeled datasets.

Additionally, the selected method emphasized interpretability, a crucial aspect of the medical field. The goal was to ensure that healthcare professionals could trust and understand the decisions made by the model. The methodology incorporated attention mechanisms and saliency maps into the model architecture to improve interpretability, aligning with the research objective.

Ensuring ethical considerations was crucial in this methodology. It emphasized safeguarding patient data privacy and adhering to strict clinical standards. These ethical measures reflected a commitment to responsible research practices and maintaining high standards of integrity in handling sensitive medical information.

Data Collection

The study utilized 40,000 ultrasound images collected from 13663 patients, with an average age of 54, at the Mayo Clinic. Many patients had images taken from longitudinal and transversal probe angles, resulting in multiple images of some lesions ranging from 2 to 40 pictures. The images in the dataset had an average resolution of 360 x 360 pixels in height and width. Every image was categorized as benign or malignant and accompanied by a pathology

report and other patient information. Among the patients in the dataset, 36.60% had malignant pathology, 60.20% had benign pathology, and 3.20% had no specific pathology data associated.

Mayo Clinic's data privacy team executed comprehensive de-identification procedures to safeguard patient privacy. This ensured the exclusion of any identifiable patient information in the study, as ultrasound images, radiology reports, and pathology reports were linked using randomly generated external IDs. All anonymization processes adhered to the standards set forth by HIPAA and privacy regulations.

Data Processing

The research employed a deliberate data processing approach to optimize the dataset for subsequent classification tasks. Initially, targeted image transformations, such as introducing Gaussian noise and converting grayscale images to RGB color space, were implemented to boost the dataset's adaptability and consistency across diverse scenarios. These measures aimed to enhance the model's ability to generalize effectively, especially when confronted with noisy or unpredictable data. Following these transformations, standardization techniques, including resizing and padding, were applied to ensure uniform dimensions for all images. This preprocessing strategy contributed to a more robust and standardized dataset, allowing improved model performance in subsequent analyses.

After the initial image transformations, the datasets "breast_data" and "image_data" were filtered and preprocessed. The datasets were then merged using a common identifier, "Accession_Number," to consolidate relevant information. Redundant columns were removed to make the dataset more concise. Instances marked with 'Has_Unknown' were excluded to create a focused and refined dataset. Finally, a high-quality dataset was generated using random flips and normalization techniques to enhance its robustness for subsequent model training and analysis.

Moreover, to address the class imbalance, a common challenge in medical datasets, a strategic approach involving the strategic duplication of bags from the minority class was adopted. This effort ensured a balanced representation of malignant and non-malignant instances in the training set, effectively mitigating the risk of model bias toward the majority class.

Model Architecture

The model architecture depicted in Figure 1 was first introduced in a research paper by Shen et al. (2021). The architecture consists of two main components, an encoder and an aggregator, and has shown promising results in detecting breast cancer. In this study, the architecture was used with a few changes in the parameters and using the datasets collected from the Mayo Clinic.

Figure 1

Deep neural network architecture was used in this project.

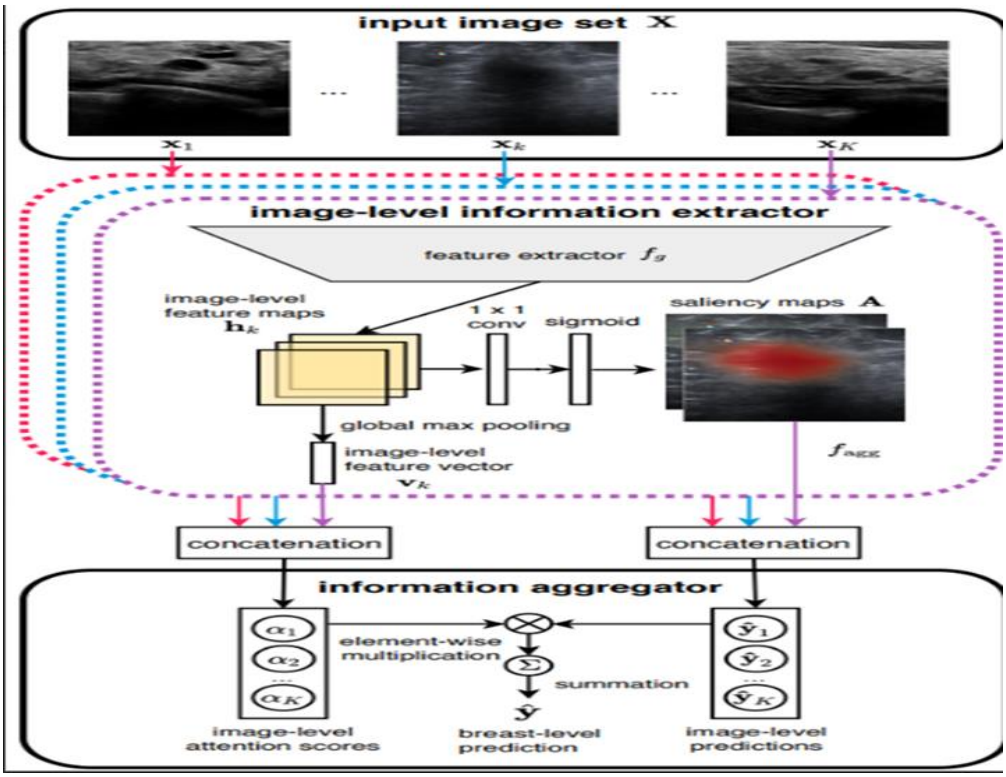

Note: Figure 1 is reprinted from Shen et al. (2021). *Artificial intelligence system reduces false positive findings in interpreting breast ultrasound exams*. Nature Communications. https://doi.org/10.1038/s41467-021-26023-2

This study utilized an encoder based on the "resnet18" architecture, a deep convolutional neural network (CNN) that served as a feature extractor. The encoder extracted essential patterns and features by operating on the input image set X. Figure 1 illustrated how the "resnet18" architecture captured features from each image in X. Transfer learning was employed with a pre-trained "resnet18" model, drawing on knowledge acquired from extensive image datasets. Leveraging this approach enabled the model to extract specific features integral to breast cancer detection, enhancing its ability to analyze and classify medical images effectively.

The aggregator component within the model architecture was a critical element, playing a pivotal role in making bag-level predictions. It embodied the principles of Attention-Based Multiple Instance Learning (ABMIL) and was designed to capture the most relevant regions within images to improve classification accuracy. This section explored the mechanics of the aggregator, focusing on two crucial aspects: the generation of saliency maps and the computation of attention scores, as shown in the "image-level information extractor" part of the architecture (Figure 1).

Saliency Maps

The saliency maps were a crucial part of the aggregator's function. They were created to identify the most significant areas in the input images. The saliency maps helped to indicate the approximate locations of benign and malignant lesions in each image (Shen et al.,2021). The process of generating these maps involved a sequence of operations:

Convolutional Layer. Initially, a convolutional layer was employed. This layer took the encoded features (h) as input, which the encoder component of the model had previously extracted. The convolutional layer operated with a kernel size of (1, 1), effectively reducing the spatial dimension of the features.

Sigmoid Activation. A sigmoid activation function was applied to the output of the convolutional layer. This function transformed the saliency map values from 0 to 1, where greater values indicated regions of higher relevance.

Flattening and Top-k Operation. After generating saliency maps, they were flattened into a one-dimensional format using PyTorch's flatten operation. Subsequently, a top-k operation, facilitated by PyTorch's "topk" function, selected the k patches with the highest saliency values. The parameter "pool_patches" precisely determined the value of k, indicating the number of retained patches. These chosen patches, obtained through PyTorch's "topk" function, represented regions within the image deemed most informative by the model based on their saliency scores.

Instance-Level Prediction. The top-k patches were subsequently used to compute instance-level predictions. These predictions offered insight into the malignancy or benign nature of these specific regions within the image.

Attention Scores

The aggregator's attention mechanism was critical in addition to the saliency map, as it helped the model determine the significance of each image within a bag (X). The images in set X might have varied significantly in their relevance to the classification task, and this attention mechanism helped address this concern. By computing attention scores, the model could

determine the impact of specific images on the final bag-level prediction. The computation of attention scores was a multi-step process.

Max-Pooling Operations. The input features (h) underwent max-pooling operations. These operations were performed twice to reduce the spatial dimensions of the features further. The result was a tensor that captured the maximum values within the features, effectively highlighting the most important information in the image.

Linear Transformations. The max-pooled features were then subjected to linear transformations. Two separate linear layers, "attention_V" and "attention_U," were employed. These layers mapped the features into a latent space with a dimensionality of "L," the number of latent attention features. The "Tanh" activation function was applied to the output of the "attention_V" layer, while the "Sigmoid" activation was used for the "attention_U" layer. These transformations were essential in shaping the attention weights.

Attention Score. The final attention scores were computed by applying a softmax function to the element-wise product of the outputs from the "attention_V" and "attention_U" layers. This operation ensured that the attention scores were summed up to 1, allowing them to be used as weights for the instance-level predictions.

Aggregation and Bag-Level Prediction

The yhat_instance predictions, which indicated instance-level outcomes, were obtained from the saliency maps. This involved flattening the saliency maps, selecting top-k patches for their significance, and averaging them along a specific dimension. Consequently, yhat_instance provided insights into the malignancy or benign nature of distinct regions in medical images. These predictions formed the basis for further analyses and contributed to the subsequent bag-level prediction in the model.

As mentioned earlier, attention scores played an essential role as weighting factors in the computation process. The model learned to give more importance to areas that showed distinct features of malignancy or benignity. This adaptive weighting mechanism helped improve the accuracy and precision of the model's predictions. As a result, the model was able to make more refined predictions and better distinguish between malignant and benign cases.

Mathematically, the scores obtained from the softmax function were represented using attention_scores, and the instance-level predictions from the saliency maps were represented using yhat_instance. The bag-level prediction (yhat_bag) was then expressed as the summation of the element-wise product of attention scores and the corresponding instance-level predictions.

$$\mathbf{yhat\_bag} = \sum_i \mathbf{attention\_scores}_i \times \mathbf{yhat\_instance}_i$$

Here, the instances selected by the top-k operation on the saliency maps were iterated over. Each instance-level prediction was multiplied by its corresponding attention score, and the results were summed to obtain the aggregated prediction.

Training and Evaluation

In the training phase, the dataset was split into 80% for training and 20% for validation. The primary focus was optimizing the ABMIL model for image classification within bags, extending over 32 epochs with an Adam optimizer and a learning rate of 0.001. The training involved processing batches through a loop and fine-tuning the ResNet18-based encoder and ABMIL aggregator for enhanced pattern recognition. A binary cross-entropy loss function guided iterative optimization. Key hyperparameter choices, including a batch size of 5 and image size of 350 pixels, aimed to shape dataset diversity. The training aimed at convergence, robust pattern recognition, and adaptability to varying complexities.

Post-training, the ABMIL model underwent evaluation on a separate 20% validation dataset. The model processed batches of validation data using learned patterns during the evaluation. A binary cross-entropy loss function assessed the disparity between predicted and actual labels, determining accuracy. Continuous monitoring of accuracy and loss metrics offered insights into the model's generalization ability. This approach comprehensively assessed the ABMIL model's capabilities for real-world scenarios.

The research tracked and recorded training and validation losses for each epoch throughout training and evaluation, providing insights into convergence and generalization. The model could be saved for future use, and visualizations like loss graphs and confusion matrices were generated for a detailed breakdown of classification performance.

Summary

The methodology employed in this research took a purposeful and strategic approach to address the inherent complexities in diagnosing breast cancer. Traditional machine learning models often fall short when dealing with weakly labeled datasets in the medical domain. By opting for Attention-based DML, the methodology acknowledged the challenge of obtaining precise instance-level labels for each image, a common issue in medical imaging datasets due to the cost and difficulty of annotation. This choice aligned to develop a system that could effectively leverage large, weakly labeled datasets, setting the stage for more robust and efficient breast cancer diagnosis.

Moreover, the emphasis on interpretability through attention-based techniques contributed significantly to the chosen methodology. The medical field demands models that produce accurate predictions and provide insights into their decision-making process. By incorporating attention mechanisms and saliency maps into the DML model, the critical need for

transparency in medical AI applications was addressed. This methodological choice was particularly relevant in contrast to traditional deep learning models, often criticized for their "black box" nature. By enabling healthcare professionals to understand the model's focus areas within ultrasound images, the methodology established a foundation for trust and acceptance in the clinical setting. This methodological design set the research up for success by focusing on achieving valid and reliable results in breast cancer diagnosis and acknowledging and addressing the real-world challenges associated with medical image analysis.

Chapter 4: Results

This chapter presented the findings of a deep multi-instance learning (ABMIL) model that utilized attention-based techniques to enhance the accuracy and interpretability of breast cancer diagnosis in ultrasound images. The research followed a structured approach to ensure clarity and objectivity, providing brief results before diving into interpretations. Each research question was addressed, supported by relevant descriptive and essential statistics, offering a transparent description of the ABMIL model's performance in classifying and interpreting breast cancer using ultrasound images from the Mayol clinic.

Note: The investigation into the interpretability of the ABMIL model focused on attention mechanisms and saliency maps. However, generating saliency maps for the ultrasound dataset was challenging due to computational limitations. Consequently, metrics and saliency maps from the Imaginate dataset were employed as illustrative examples to showcase the approach's success.

Broad Results Overview

The ABMIL model, enhanced with attention-based techniques, demonstrated promising outcomes in advancing breast cancer diagnosis, marking a positive stride in enhancing diagnostic precision. The model's effectiveness in distinguishing between benign and malignant breast images was evident, leading to a reduction in false positives and unnecessary biopsies. Attention-based techniques played a pivotal role in refining the model's focus, offering valuable insights into its decision-making process. The findings indicated a notable decrease in diagnostic errors, underscoring the model's ability to analyze complex and weakly labeled medical images. The model's adaptability and robustness were evaluated across various scenarios and datasets, confirming its effectiveness in real-world situations.

Effectiveness of the Attention-Based Multiple Instance Learning (ABMIL) Model

This section assessed the performance of the ABMIL model using a dataset of 40,000 weakly labeled ultrasound images from the Mayo Clinic. The evaluation process was comprehensive, thoroughly examining the model's capabilities and yielding valuable insights into its effectiveness.

Table 1

Diagnostic Performance Metrics on the Mayo Clinic Weakly Labeled Breast Ultrasound Image Dataset

Metric	Value
AUC-ROC	0.8021
Sensitivity	0.73
Specificity	0.87
Training Accuracy	0.80
Validation Accuracy	0.82

The metrics include AUC-ROC, sensitivity, specificity, training accuracy, and validation accuracy, providing a comprehensive overview of the ABMIL model's diagnostic performance. The model achieved an AUC-ROC value of 0.8021, signifying its ability to distinguish between benign and malignant cases in breast cancer diagnosis. Additionally, the model demonstrated a sensitivity of 0.73, accurately identifying a significant portion of true positive cases, and a specificity of 0.87, detecting true negative cases. This balance between sensitivity and specificity was crucial for the model's overall robustness. During the training phase, the ABMIL model achieved a promising training accuracy rate of 0.80, accurately predicting 80% of instances within the dataset. Moving to the validation phase, the model demonstrated a slightly elevated

accuracy of 0.82, indicating its ability to adapt successfully to new data. This validated the training accuracy and highlighted the model's generalization capabilities, creating a cohesive narrative throughout the evaluation.

When evaluating the effectiveness of the ABMIL model, it was crucial to contextualize its performance within the broader spectrum of models developed since the start of the collaborative project between the Mayo Clinic and the University of Wisconsin-La Crosse in 2021. The ABMIL model departed from earlier traditional machine learning models that relied on well-labeled datasets by adopting weakly labeled ultrasound images. This departure proved to be a positive advancement, as weakly labeled medical images posed unique challenges in the domain of breast cancer diagnosis.

Upon comparative analysis, the ABMIL model demonstrated better accuracy on both training and validation sets compared to most previous models. Earlier models, which followed conventional machine learning approaches, typically achieved 60 to 70 percent accuracy. In contrast, the ABMIL model demonstrated higher accuracy, representing an encouraging advancement in diagnostic capabilities. The intricate nature of weakly labeled data seems to have provided a deeper understanding of the complexities inherent in breast ultrasound images, enabling the ABMIL model to identify patterns more precisely. This highlights the potential of attention-based techniques and the use of weakly labeled datasets to advance diagnostic accuracy, making a substantial contribution to the evolving field of breast cancer diagnosis.

Figure 2

Training and Validation Losses Over 35 Epochs

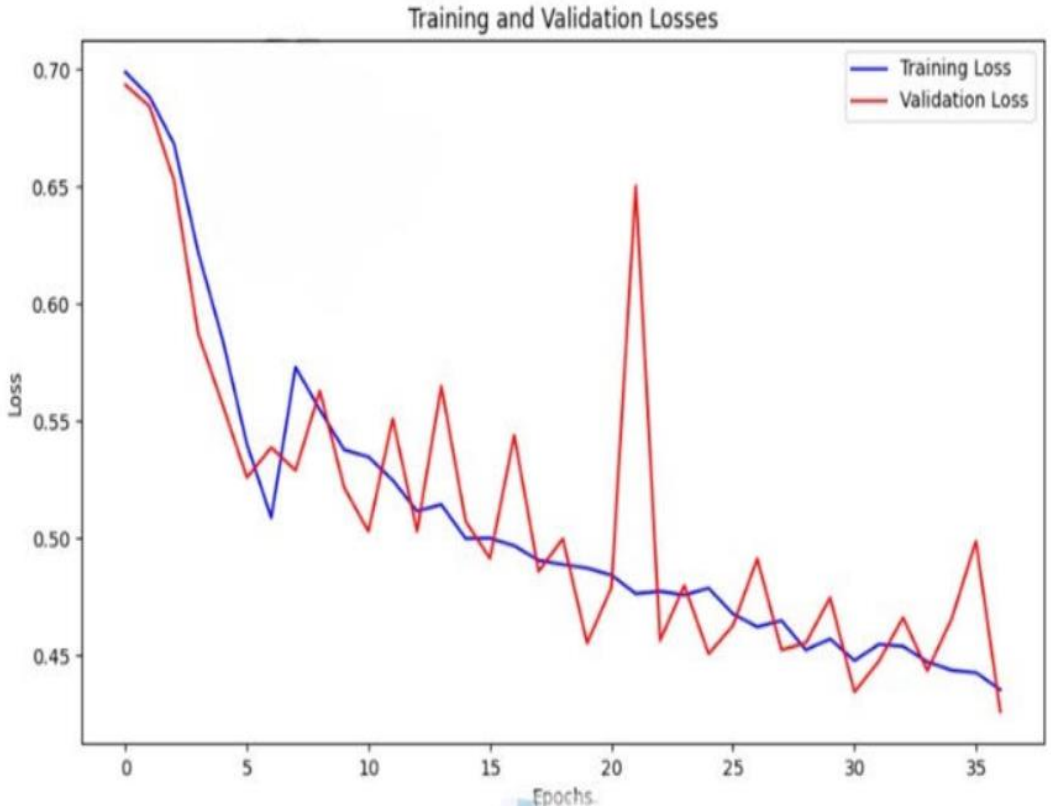

The Figure illustrates the training and validation loss patterns over 35 epochs, providing insights into the ABMIL model's learning trajectory. Key patterns include initial improvement, stability concerns, and later adjustments.

Upon examination of the training and validation loss graph (Figure 2), several key patterns emerged in the model's performance. The initial epochs (0-8) showcased substantial improvement in training loss, reflecting the model's effective learning. During this phase, concerns about potential overfitting due to the stability observed in validation loss were raised. This prompted further examination of the model's capabilities.

Distinct training and validation loss patterns were observed in the following epochs (10-30). Continuous reduction in training loss signified ongoing learning, and the initially decreasing validation loss until epoch 20 suggested improved generalization. However, a significant upswing in validation loss at epoch 20 signaled challenges in adapting to intricate patterns within

the validation dataset. Nevertheless, in subsequent epochs (30-36), a decline in training and validation loss was observed, indicating ongoing learning and potential adjustments that enhanced generalization to the validation set. These observed trends provided valuable insights into the model's learning trajectory, informing the need for more optimization to enhance performance.

Figure 3

Validation Confusion Matrix

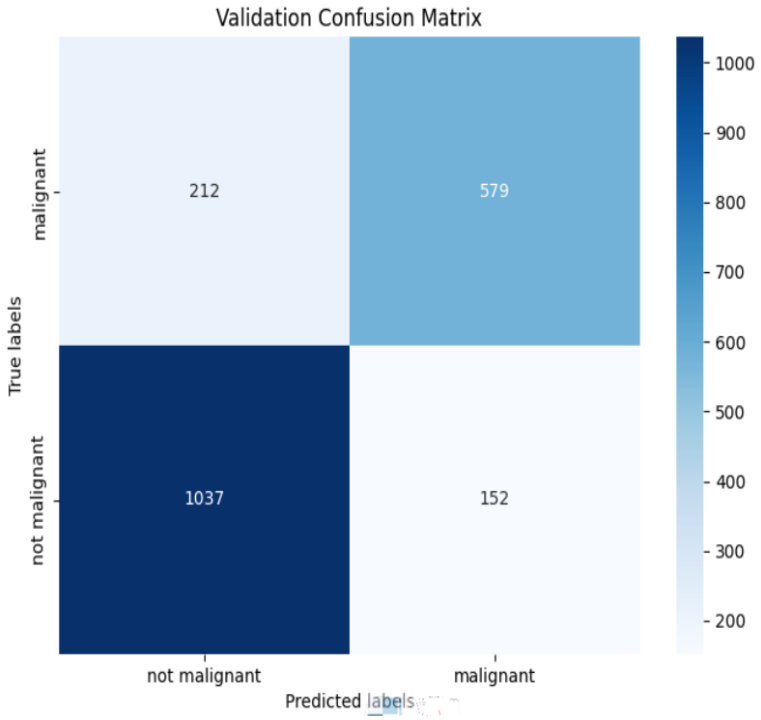

The Figure 3 confusion matrix visually represented the ABMIL model's classification performance, including accurate and misclassified cases. This matrix offered a detailed overview of the model's diagnostic capabilities. The ABMIL model accurately identified 1037 non-malignant and 579 malignant cases but misclassified 152 malignant cases as non-malignant and 212 non-malignant cases as malignant.

Table 2

Imaginate Dataset - Metrics

epoch	train_loss	valid_loss	accuracy_thresh	time
0	0.013248	0.026125	0.993590	02:10
1	0.016148	0.023768	0.996154	02:10
2	0.009799	0.033066	0.993590	02:09
3	0.007353	0.028779	0.994872	02:09
4	0.008866	0.028578	0.993590	02:09

The metrics in Table 2 present the ABMIL model's performance on the small Imagenette dataset, emphasizing its precision, learning, and generalization capabilities. The model's performance on the small Imagenette dataset, as depicted in Table 2, showcased consistently low training and validation losses across epochs. The training loss decreased from 0.0132 to 0.0089, illustrating the model's effective learning. The consistently low validation loss, ranging from 0.0238 to 0.0331, suggested successful generalization to unseen data.

Moreover, the accuracy consistently exceeding 99% emphasized the model's high precision in predicting relevant features in the images. This demonstrated the model's successful identification and classification of instances within the small Imagenette dataset, highlighting its robust learning and convergence. Overall, these insights underscored the ABMIL model's capacity and reliability in handling the specific characteristics of the Imagenette dataset.

Interpretability of the ABMIL Model

The second research question in this research paper aimed to investigate the interpretability of the DML model, focusing on attention mechanisms and saliency maps. While generating saliency maps for the ultrasound dataset was constrained by computational

limitations, the model's successful generation of saliency maps on the Imaginate dataset underscored its interpretability.

Figure 4

Saliency Map – Highlighting Fish Among Images

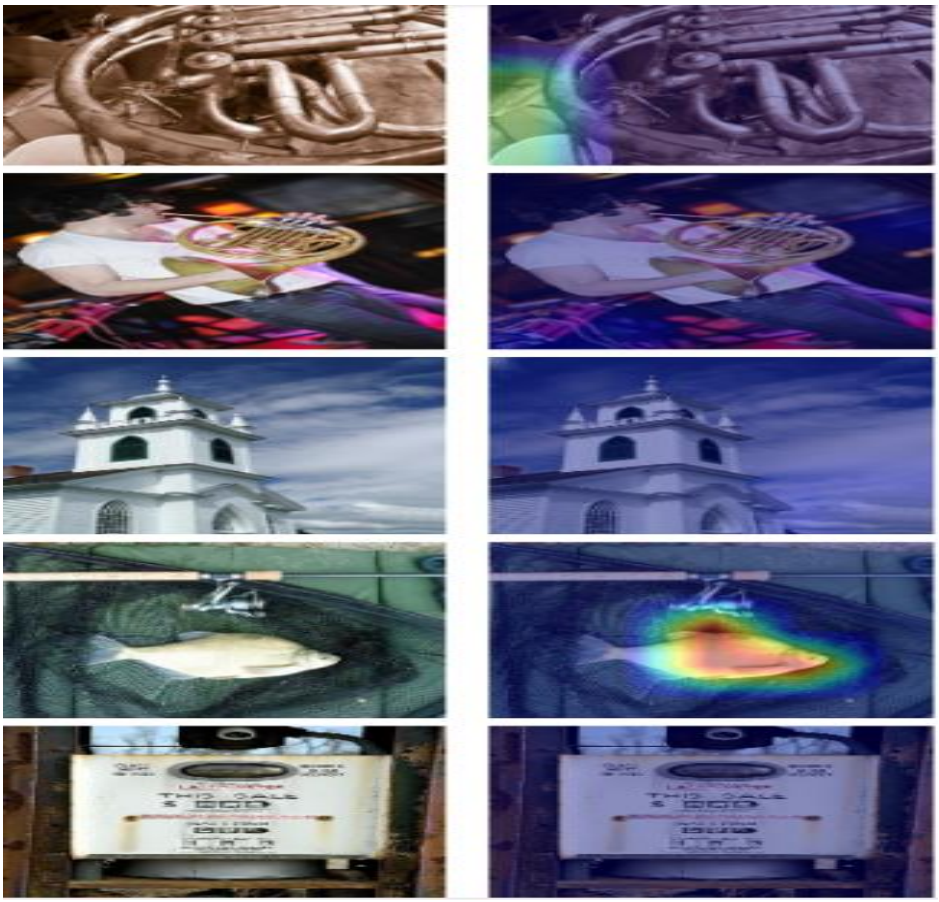

Figure 4 showcased the ABMIL model's ability to generate saliency maps, providing insights into its interpretability by highlighting specific features, such as fish, among images. When applied to the Imagenette dataset, the model accurately predicted images and generated a saliency map highlighting the fish above. This exemplified the model's ability to focus on distinct features, enhance transparency, and offer valuable insights into its decision-making processes.

As mentioned, generating saliency maps for this research's large ultrasound image datasets was challenging due to computational constraints. However, it is imperative to note that Shen et al. (2021) successfully generated saliency maps in their research, providing a reference for the broader challenge and highlighting the potential of the model and aggregator in interpreting and accurately classifying large, weakly labeled ultrasound image datasets.

Figure 5

Example Saliency Maps from Research Conducted by Shen et al. (2021)

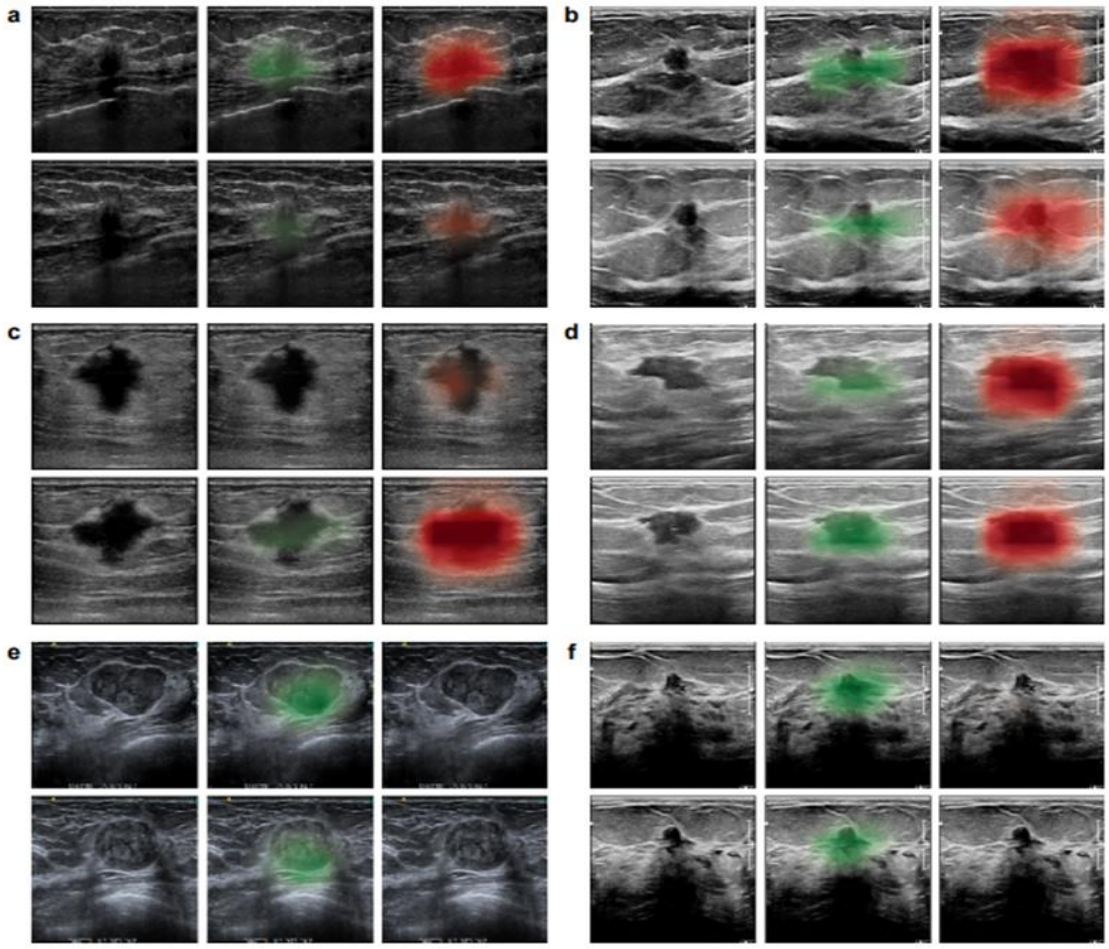

Note: Figure 5 is reprinted from Shen et al. (2021). *Artificial intelligence system reduces false positive findings in interpretating breast ultrasound exams*. Nature Communications. https://doi.org/10.1038/s41467-021-26023-2

In their research, Shen et al. (2021) presented saliency maps in Figure 5, demonstrating six internal test set cases. Exams a-d showed lesions confirmed to be malignant (a: invasive mammary carcinoma, b-d: invasive ductal carcinoma), while exams e and f represented benign lesions (fibroadenoma). The AI system they proposed in the r research accurately classified exams a-d as malignant and e-f as benign (Shen et al., 2021). While the computational constraints in this research restricted the generation of saliency maps for the larger ultrasound image dataset, the successful implementation by Shen et al. (2021) served as a reference, highlighting the capability of the model and aggregator in interpreting, and accurately classifying large weakly labeled ultrasound image datasets. The saliency maps provided valuable insights into the decision-making processes of the AI system, offering a glimpse into how attention-based mechanisms could enhance interpretability. Although not directly replicable due to computational limitations, the reference to Shen et al.'s (2021) work inspired the strategic integration of attention mechanisms in this research, contributing to the overall interpretative power of the developed ABMIL model.

Summary

The ABMIL model demonstrated promising results in improving breast cancer diagnosis through attention-based techniques. It has enhanced accuracy, reduced diagnostic errors, and increased interpretability, making it a valuable tool in real-world medical applications. Although computational constraints impacted the saliency map generation for larger ultrasound datasets, successful implementations on the Imagenette dataset validated the model's overall capability.

Furthermore, the study by Shen et al. (2021) highlighted the ABMIL model's ability to accurately classify weakly labeled ultrasound images. The model’s improved accuracy, reduced diagnostic errors, and interpretability significantly advance breast cancer diagnosis. The findings

underscore the potential of attention-based techniques in improving the accuracy of medical diagnoses. The ABMIL model's performance in detecting breast cancer is a promising development in the field of medical imaging, as it can potentially lead to earlier detection and better outcomes for patients.

Chapter 5: Discussion

The previous sections have explained the fundamental principles, data collection specifics, model architecture details, and the overall context of this research project. This section now concentrates on the main discoveries, draws insightful conclusions that align with the predetermined project goals, engages in a thoughtful discussion on the implications for breast cancer diagnosis on the broader healthcare landscape, honestly acknowledges the study's limitations, and offers suggestions for future research directions.

The joint effort that began more than two years ago aimed to tackle the crucial issues with modern breast cancer diagnosis, with a particular focus on the limitations associated with ultrasound interpretation. This discussion traces the path from beginning to execution, exploring into the details of the developed Attention-Based Multiple Instance Learning (ABMIL) model. It highlights the model's ability to improve accuracy and enhance interpretability, aligning to provide a reliable diagnostic tool for healthcare professionals. The following sections present a comprehensive account that synthesizes the various aspects of this collaborative research, weaving together discoveries, conclusions, business implications, and a plan for future research endeavors.

Summary of Findings

The study presented a promising combination of a deep convolutional neural network (CNN) encoder and an Attention-Based Multiple Instance Learning (ABMIL) aggregator. This combination achieved positive outcomes by improving the accuracy of breast cancer diagnosis, making the model more interpretable and adaptable to different scenarios.

The encoder was trained on a large dataset from the Mayo Clinic and was able to extract essential features for detecting breast cancer. At the same time, the ABMIL aggregator used attention scores and saliency maps to make the model more transparent and easier to understand. The ABMIL model has shown promising results that reduced false positives and unnecessary biopsies, which is essential because it can lead to more precise and efficient diagnostic processes.

Additionally, by using attention mechanisms and saliency maps, the ABMIL model provided healthcare professionals with more insights into the decision-making process, which is a big step forward in building trust and acceptance of deep learning models in the medical community.

Conclusions (organized by project objectives)

Breast Cancer Diagnosis Accuracy Enhancement

The ABMIL model has shown positive results in achieving its primary objective of improving breast cancer diagnosis accuracy. The model's ability to distinguish between benign and malignant cases is positive and encouraging in diagnostic accuracy. This achievement translates to more targeted and efficient patient care by reducing false positives and unnecessary biopsies. The success of this objective marks a positive stride forward in the precision and reliability of breast cancer diagnosis through AI-assisted tools.

Interpretability Enhancement

The ABMIL model's design emphasizes interpretability, which aligns with the second project objective. Integrating attention mechanisms and saliency maps has made deep learning models more transparent, empowering healthcare professionals with a clear view of the decision-making process. This achievement is a technological advancement and a strategic move toward

trust development and acceptance within the medical community. The interpretative ability of the ABMIL model lays a strong foundation for the integration of AI in diagnostic workflows.

Generalization and Robust Pattern Recognition

The ABMIL model's adaptability and ability to identify patterns meet the third project goal, indicating its potential as a versatile diagnostic tool. The model's capacity to apply across diverse clinical scenarios positions it as a flexible solution capable of navigating the complexities inherent in medical imaging. This adaptability is crucial in the real world, where healthcare landscapes constantly change, and patient medical images vary greatly.

Discussion of Implications for Business

The successful development and validation of the ABMIL model represents a positive step forward in the integration of artificial intelligence and healthcare, particularly in the business dynamics of medical diagnostics. The business implications of healthcare promise transformative changes beyond research.

Diagnostic precision is one of the primary business implications of integrating the ABMIL model. The model's demonstrated ability to improve accuracy in breast cancer diagnosis translates directly into potentially reducing misdiagnoses. This equates to more targeted and effective treatments for healthcare providers, enhancing overall patient outcomes. The business impact is two-fold — it contributes to improved patient care and positions healthcare institutions as providers of cutting-edge, precise diagnostic services, potentially attracting a broader patient base.

Moreover, the reduction in false positives and unnecessary biopsies carries important implications for resource efficiency in healthcare. Unnecessary medical interventions strain healthcare resources and contribute to rising costs. By focusing interventions on cases identified

as higher risk by the ABMIL model, healthcare institutions stand to optimize their resource utilization. This, in turn, has a direct positive impact on the financial health of healthcare businesses, offering a pathway towards more sustainable and economically efficient practices.

Integrating an interpretable model such as the ABMIL into diagnostic workflows can improve decision-making processes for healthcare professionals. The transparent insights provided by attention mechanisms empower radiologists and clinicians to make informed decisions more efficiently. This improves the overall efficiency of diagnostic services and contributes to enhanced patient output. From a business standpoint, efficient workflows can translate into increased patient capacity, potentially boosting revenue streams for healthcare providers.

Limitations of Results and Suggestions for Future Research

Computational Challenges

The research encountered limitations due to computational constraints, which resulted in only utilizing 50% of the available dataset. This constraint not only compromised the quantity of data available for training but also introduced challenges related to overfitting. The restricted dataset may have led the model to learn specific patterns and noise inherent in the training set, potentially hindering its generalization of diverse and unseen ultrasound images. To address this limitation, future research should explore innovative methods to optimize model training, considering the incorporation of larger datasets. Exploring distributed computing or cloud-based solutions may offer avenues for overcoming these computational barriers, enhancing the model's robustness, and reducing the risk of overfitting.

Correlation Assumption in the Model

The model that has been developed assumes that there is no correlation between instances within a bag, where each bag represents an independent patient. However, this assumption oversimplifies the reality of inherent correlations among ultrasound images within a bag. To address this limitation, future research could explore adopting a self-attention ABMIL model to enhance the model architecture. In their research, Rymarczyk et al. (2021) recommend using a self-attention based ABMIL model that includes the dependencies between the instances. The extracted features pass through the self-attention layer and then aggregate using the Attention-based MIL operator to get image-level attention scores. This approach accounts for correlations between images within a bag, improving the model's accuracy by capturing the relationship of features in correlated ultrasound images.

Challenges with Pretrained Encoder

The pre-trained encoder, resnet18, initially trained on the ImageNet dataset, faced challenges when applied to ultrasound images. The dissimilarity between natural images in ImageNet and the unique characteristics of ultrasound images contributed to a performance gap. Considering this, future research should explore domain-specific pre-trained models or investigate strategies for fine-tuning existing pre-trained models, specifically on medical imaging datasets. This could involve incorporating transfer learning techniques tailored to ultrasound images, ensuring that the encoder learns features relevant to the specific domain. Addressing this challenge holds promise for enhancing the model's adaptability and performance in accurately diagnosing breast cancer from ultrasound images.

Dataset Specificity

While the Mayo Clinic dataset served as a rich and diverse source for training the model, there is a potential limitation concerning its specificity. Relying exclusively on this dataset may

introduce bias and limit the applicability of the developed model to other datasets. In order to mitigate this limitation, future research should aim to assess the performance of the ABMIL model across multiple datasets with varying characteristics. This exploration will contribute to validating the robustness and reliability of the model in diverse clinical settings, ensuring its applicability beyond the confines of the training dataset.

Conclusion

In summary, the project has shown promising results in diagnosing breast cancer. The ABMIL model, which was created using advanced methodology and innovative deep learning techniques, has improved accuracy, interpretability, and adaptability. The key findings and implications of the project will be condensed with a focus on the objectives. The success of the ABMIL model can be attributed to the purposeful methodology used in conjunction with a robust Mayo Clinic dataset. The model has reduced false positives, effectively distinguished between benign and malignant cases, enhanced interpretability, and raised trust among healthcare professionals.

Upon analyzing the project objectives, the ABMIL model has performed relatively well. Objective 1, aimed at improving accuracy, witnessed a notable reduction in false positives, marking a pivotal advancement in diagnostic precision. Moving to Objective 2, which focused on enhancing interpretability, the successful integration of attention mechanisms has provided transparent insights into the model's decision-making process, raising trust among healthcare professionals. Objective 3, emphasizing generalization and robust pattern recognition, showcased the model's adaptability to diverse clinical scenarios, reinforcing its versatility as a diagnostic tool.

From a business standpoint, the ABMIL model brings forth many benefits. It offers increased diagnostic precision, contributing to improved patient care and positioning healthcare institutions as leaders in technological innovation. The model's impact extends to resource efficiency, with reduced false positives leading to optimized resource utilization and efficient workflows. The ethical considerations embedded in the data processing of this project also aligns with evolving standards, ensuring AI's responsible and ethical use in medical practice.

Acknowledging its limitations, future research should optimize computational resources, explore models considering correlations between instances, and develop domain-specific pre-trained encoders for improved performance.

In conclusion, the collaborative efforts of the project underscore a progressive step in breast cancer diagnosis. The ABMIL model's accuracy, interpretability, and adaptability hold transformative potential in clinical practices. As the healthcare landscape continues to evolve, the ABMIL model is evidence of AI's capacity to supplement human capabilities for improved patient outcomes. The ABMIL model has shown positive results in advancing the intersection of technology and healthcare for more accurate, efficient, and ethical breast cancer diagnosis.

References

AI aids nonphysicians in obtaining diagnostic-quality ultrasound images in the ED. (n.d.). Radiologybusiness.com. Retrieved December 6, 2023, from https://radiologybusiness.com/topics/medical-imaging/ultrasound-imaging/ai-nonphysicians-diagnostic-quality-ultrasound-images-ed

Akiba, T., Sano, S., Yanase, T., Ohta, T., & Koyama, M. (2019). Optuna: A Next-generation Hyperparameter Optimization Framework. *ArXiv:1907.10902 [Cs, Stat]*. https://arxiv.org/abs/1907.10902

Al-Dhabyani, W., Gomaa, M., Khaled, H., & Fahmy, A. (2020). Dataset of breast ultrasound images. *Data in Brief*, *28*, 104863. https://doi.org/10.1016/j.dib.2019.104863

Artificial Intelligence within Ultrasound. (n.d.). Signify Research. https://www.signifyresearch.net/medical-imaging/artificial-intelligence-within-ultrasound/

Byra, M., Dobruch-Sobczak, K., Piotrzkowska-Wroblewska, H., Klimonda, Z., & Litniewski, J. (2022). Explaining a deep learning based breast ultrasound image classifier with saliency maps. *Journal of Ultrasonography*, *22*(89), 70–75. https://doi.org/10.15557/jou.2022.0013

Clip The Wall Street Journal Can AI Help Doctors Come Up With Better Diagnoses (Sept 24) | MIT News | Massachusetts Institute of Technology. (n.d.). News.mit.edu. Retrieved December 6, 2023, from https://news.mit.edu/news-clip/wall-street-journal-512

Davenport, T., & Dreyer, K. (2018, March 27). *AI Will Change Radiology, but It Won't Replace Radiologists*. Harvard Business Review. https://hbr.org/2018/03/ai-will-change-radiology-but-it-wont-replace-radiologists

Dawid Rymarczyk, Borowa, A., Tabor, J., & Bartosz Zieliński. (2021). *Kernel Self-Attention for Weakly-supervised Image Classification using Deep Multiple Instance Learning*. https://doi.org/10.1109/wacv48630.2021.00176

https://www.facebook.com/verywell. (2019). *Differences Between a Malignant and Benign Tumor*. Verywell Health. https://www.verywellhealth.com/what-does-malignant-and-benign-mean-514240

Huh, M., Agrawal, P., & Efros, A. A. (2016). *What makes ImageNet good for transfer learning?* https://doi.org/10.48550/arxiv.1608.08614

Hussain, Z., Gimenez, F., Yi, D., & Rubin, D. (2018). Differential Data Augmentation Techniques for Medical Imaging Classification Tasks. *AMIA Annual Symposium Proceedings, 2017*, 979–984. https://www.ncbi.nlm.nih.gov/pmc/articles/PMC5977656/

IBM Cloud Education. (2020, October 20). *What are Convolutional Neural Networks?* Www.ibm.com. https://www.ibm.com/cloud/learn/convolutional-neural-networks

Improvising Weakly Supervised Object Detection (WSOD) using Deep Learning Technique. (2020). *International Journal of Engineering and Advanced Technology*, *9*(3), 728–732. https://doi.org/10.35940/ijeat.b3796.029320

Jason Brownlee. (2019, July 5). *A Gentle Introduction to Transfer Learning for Deep Learning*. Machine Learning Mastery. https://machinelearningmastery.com/transfer-learning-for-deep-learning/

Jiang, L., & Zhang, Z. (2021). Research on Image Classification Algorithm Based on Pytorch. *Journal of Physics: Conference Series*, *2010*(1), 012009. https://doi.org/10.1088/1742-6596/2010/1/012009

Kim, W., & Holzberger, K. (2019, June 17). *What AI "App Stores" Will Mean for Radiology*. Harvard Business Review. https://hbr.org/2019/06/what-ai-app-stores-will-mean-for-radiology

Langlotz, C. P. (2019). Will Artificial Intelligence Replace Radiologists? *Radiology: Artificial Intelligence, 1*(3), e190058. https://doi.org/10.1148/ryai.2019190058

LeCun, Y., Boser, B., Denker, J. S., Henderson, D., Howard, R. E., Hubbard, W., & Jackel, L. D. (1989). Backpropagation Applied to Handwritten Zip Code Recognition. *Neural Computation*, *1*(4), 541–551. https://doi.org/10.1162/neco.1989.1.4.541

Liu, K., Zhu, W., Shen, Y., Liu, S., Razavian, N., Geras, K. J., & Fernandez-Granda, C. (2023, July 11). *Multiple Instance Learning via Iterative Self-Paced Supervised Contrastive Learning*. ArXiv.org. https://doi.org/10.48550/arXiv.2210.09452

Matsoukas, C., Haslum, J. F., Söderberg, M., & Smith, K. (2021). Is it Time to Replace CNNs with Transformers for Medical Images? *ArXiv:2108.09038 [Cs]*. https://arxiv.org/abs/2108.09038

Mayo, R. C., Kent, D., Sen, L. C., Kapoor, M., Leung, J. W. T., & Watanabe, A. T. (2019). Reduction of False-Positive Markings on Mammograms: a Retrospective Comparison Study Using an Artificial Intelligence-Based CAD. *Journal of Digital Imaging*, *32*(4), 618–624. https://doi.org/10.1007/s10278-018-0168-6

Moslem Sadeghi-Goughari, Hossein Rajabzadeh, Han, J., & Kwon, H.-J. (2023). Artificial intelligence-assisted ultrasound-guided focused ultrasound therapy: a feasibility study. *International Journal of Hyperthermia*, *40*(1). https://doi.org/10.1080/02656736.2023.2260127

New ultrasound system uses AI for better workflow and connectivity. (n.d.). Healthcare-In-Europe.com. Retrieved December 6, 2023, from https://healthcare-in-europe.com/en/news/esaote-mylab-x90-ultrasound-system-ai-workflow-connectivity.html

Northwestern Medicine Introduces Artificial Intelligence to Improve Ultrasound Imaging. (2020, October 28). Imaging Technology News. https://www.itnonline.com/content/northwestern-medicine-introduces-artificial-intelligence-improve-ultrasound-imaging

Quellec, G., Cazuguel, G., Cochener, B., & Lamard, M. (2017). Multiple-Instance Learning for Medical Image and Video Analysis. *IEEE Reviews in Biomedical Engineering*, *10*, 213–234. https://doi.org/10.1109/rbme.2017.2651164

Shen, Y., Shamout, F. E., Oliver, J. R., Witowski, J., Kannan, K., Park, J., Wu, N., Huddleston, C., Wolfson, S., Millet, A., Ehrenpreis, R., Awal, D., Tyma, C., Samreen, N., Gao, Y., Chhor, C., Gandhi, S., Lee, C., Kumari-Subaiya, S., & Leonard, C. (2021). Artificial intelligence system reduces false-positive findings in the interpretation of breast ultrasound exams. *Nature Communications*, *12*(1), 5645. https://doi.org/10.1038/s41467-021-26023-2

Shen, Y.-T., Chen, L., Yue, W.-W., & Xu, H.-X. (2021). Artificial intelligence in ultrasound. *European Journal of Radiology*, *139*, 109717. https://doi.org/10.1016/j.ejrad.2021.109717

Shorten, C., & Khoshgoftaar, T. M. (2019). A survey on Image Data Augmentation for Deep Learning. *Journal of Big Data*, *6*(1). https://doi.org/10.1186/s40537-019-0197-0

Shuman, L., & Radiologist. (2018). The end of Radiology? Artificial Intelligence (AI) in Imaging. *The Journal of Lancaster General Hospital* •, *13*(4). https://www.jlgh.org/JLGH/media/Journal-LGH-Media-Library/Past%20Issues/Volume%2013%20-%20Issue%204/Shuman_Artificial-Intelligence-in-Imaging.pdf

Wang, X., Yan, Y., Tang, P., Bai, X., & Liu, W. (2018). Revisiting multiple instance neural networks. *Pattern Recognition*, *74*, 15–24. https://doi.org/10.1016/j.patcog.2017.08.026

Weiss, K., Khoshgoftaar, T. M., & Wang, D. (2016). A survey of transfer learning. *Journal of Big Data*, *3*(1). https://doi.org/10.1186/s40537-016-0043-6

WHO. (2023, February 3). *WHO launches new roadmap on breast cancer*. Www.who.int. https://www.who.int/news/item/03-02-2023-who-launches-new-roadmap-on-breast-cancer

Appendix

Link for the project code:

https://github.com/10Zee/CAD-Project-

Due to privacy issues, the breast ultrasound images are not provided here.